

Architecture — CareBot Catch & Cage (Theater)

Diagrams for facilitators and for the technical audiences who ask "what is actually happening here?"

Every diagram is an **Excalidraw** drawing. Each one ships twice: an **.svg** that renders here, and an **.excalidraw** source you can open at excalidraw.com to edit, rebrand or lift a single panel into a slide. Nothing needs to be installed to do that.


Read this first. The frontend is a single self-contained **index.html** with no backend and no dependencies. The Microsoft pipeline it depicts is **simulated in the browser**. The alert catalog, MITRE mappings and Graph operations are transcribed from Microsoft's published documentation so the depiction is accurate, but nothing here calls a live tenant. See what is real versus simulated.

1. The story: attack, detect, contain

The full arc the audience watches, end to end.

Attack → Detect → Contain

What the audience watches, end to end



Detection surfaces ~15s after the prompt. Containment then runs 18-22 steps at ~0.6s each, so the agent is fully caged about 26-28 seconds after the attack. Both are tunable in Settings.

Attack, detect, contain

Edit: 01-attack-detect-contain.excalidraw

Timing the audience sees: detection surfaces about 15 seconds after the prompt, which mirrors real SOC ingest and analytics latency. Containment then runs 18 to 22 steps at roughly 0.6 seconds each, so the agent is fully caged about 26 to 28 seconds after the attack. Both are tunable in Settings.


Two scoring points, not one. The prompt is scored on the way in, before the model answers, because that is where detection actually happens. The response is scored on the way out, which is what confirms a compromise. Saying this out loud lands well with practitioners.

2. The layered cage

Containment is not one action. It scales with what the attacker actually did, which is the part security practitioners care about most.

The layered cage

Containment scales with what the attacker actually did



Network is added for exfiltration, C2 or High severity. The compute layer is the patient-safety kill switch and is gated on purpose: a real SOC puts an approval in front of that blast radius, so it sits behind a Settings toggle. Turning it off is a good way to show graduated response.

The layered cage

Edit: 02-layered-cage.excalidraw

Why identity comes first. An agent is a workload identity. Disabling it is the decisive control, and it is the one step that is genuinely verified rather than illustrative. Data revocation is the backstop for the window where an already-issued token is still technically valid.

Why the last layer is gated. The compute and model hard stop is the patient-safety kill switch. A real SOC would put an approval in front of that blast radius, so it sits behind a Settings toggle. Turning it off is a good way to show graduated response.


On continuous access evaluation. CAE covers single-tenant service principals, requires Workload Identities Premium, and does not support managed identities. Where it does not apply, an issued access token stays valid until it expires, which is exactly why the data-plane revocations matter. The demo says this on screen rather than implying instant revocation everywhere.

3. What is real versus simulated

The most important diagram in this file. Use it when a technical audience asks what they are actually looking at.

What is real, what is simulated

The question every technical audience asks



"The detection pipeline is simulated so this runs anywhere with no tenant and no cost. What is not simulated is the shape of it. Every alert ID, severity and tactic on that screen is transcribed from Microsoft's published catalog, and the containment steps are the real Graph and Azure operations a playbook would call."

What is real, what is simulated

Edit: 03-real-vs-simulated.excalidraw

The honest line to use out loud: "The detection pipeline is simulated so this runs anywhere with no tenant and no cost. What is not simulated is the shape of it. Every alert ID, severity and tactic on that screen is transcribed from Microsoft's published catalog, and the containment steps are the real Graph and Azure operations a playbook would call."

Alerts, pinned to the published catalog

Alert metadata is a property of the alert, not of the attack that tripped it, so the same ID always carries the same tactics and severity.

Attack in the demo	Defender alert	ATT&CK tactics	Severity
Credential theft	AI.Azure_CredentialTheftAttempt	Credential Access, Lateral Movement, Exfiltration	Medium
ASCII smuggling	AI.Azure_ASCIIsmuggling	Impact	High
Unsafe action, prior-auth tampering	AI.Azure_AnomalousToolInvocation	Execution	Low
Jailbreak, PHI exfiltration, prompt leak	AI.Azure_Jailbreak.ContentFiltering.DetectedAttempt	Privilege Escalation, Defense Evasion	Medium

Source: [Alerts for AI services](#).

A teaching moment worth using. The tool-invocation alert is genuinely **Low**, yet it still drives a **High** incident. That is exactly how a real SOC works: you triage the incident, not the raw alert.


Product status. Defender for Cloud threat protection for AI is generally available. Entra ID Protection for Agents is newer and licensing-gated, so it appears as an illustrative step rather than a live tenant call.

4. Three-screen Theater wall

One PC, three browser windows, three displays. The panels stay in sync without a server.

Three-screen Theater wall

48 Ft x 9 Ft canvas · three 16:9 tiles · one media PC



Panel ① is the brain. Only the Agent panel drives state; ② and ③ mirror it from a broadcast snapshot, so they can never disagree with what the audience just did.

⚠ Sync needs ONE shared HTTPS origin. On file:// paths each window is a separate opaque origin and the panels silently stop syncing. Always launch from the hosted URL.

Three-screen Theater wall

Edit: 04-theater-wall.excalidraw

Panel ① is the brain. Only the Agent panel drives state. Panels ② and ③ mirror it from a broadcast snapshot, which is why they can never disagree with what the audience just did.

Why a shared origin matters. Sync uses `BroadcastChannel` and `localStorage`, which require one shared HTTPS origin. Launch all three panels from the same deployed hostname. On `file://` paths each window is a separate opaque origin and sync silently breaks.

A detail worth pointing at. The wall exposes the detection dwell gap: Panel ① turns red the moment the agent is compromised while Panel ③ still reads ACTIVE, until containment catches up. Security practitioners notice.


Launch it: open `?panel=wall` or click **Wall · 3-screen**, then send each panel to its display.

5. Hosting

Three ways to run it, in increasing order of control.

Hosting

Three ways to run it, in increasing order of control



Do not gate a walk-up floor. A tenant gate blocks anonymous phone joins, so the QR experience stops working for anyone outside your tenant. Gate it only for internal-only audiences.

Hosting options

Edit: 05-hosting.excalidraw

Path	Use it when	Cost
GitHub Pages	You just want to run the demo	\$0
Hub PC shortcuts	A dedicated kiosk or wall PC, with an offline backup copy	\$0
Azure SWA, anonymous	You want your own hostname and URL control	\$0, Free plan
Azure SWA, Entra gated	Internal-only audiences	About \$9/mo, Standard plan


Choosing the gate. A tenant gate blocks anonymous phone joins, so the QR experience stops working for anyone outside your tenant. For a walk-up Hub floor, stay anonymous. Detail in [azure-deploy-scripts.md](#).

6. Inside the frontend

For anyone forking it.

Inside the frontend

One file, for anyone forking it



Scoring happens on the way IN, before the model answers, because detection is on the inbound prompt. Output-side checks then confirm a compromise. Mock is deterministic and offline, which is what you want in front of an audience.

Inside the frontend

Edit: 06-frontend-internals.excalidraw

- **Mock engine** is deterministic and offline. Same input, same output, every time. This is what you want in front of an audience.
- **Live engine** calls an Azure OpenAI or OpenAI-compatible endpoint. The key is stored only in that browser and is never committed.
- **Scoring happens on the way in**, before the model answers, because detection is on the inbound prompt. Output-side checks then confirm a compromise.

Editing and reusing these diagrams

The `.excalidraw` files in diagrams/ are the source.

1. Go to excalidraw.com.
2. **File** → **Open**, pick the `.excalidraw` file.
3. Edit freely. Rebrand it, drop a panel, translate the labels.
4. **File** → **Export image** → **SVG** back over the matching `.svg` if you want the copy in this repo to change too.

Lifting one into a slide: select the elements you want, right-click and **Copy to clipboard as PNG**, then paste straight into PowerPoint.

The palette matches the app: blue `#0f6cbd` for the Microsoft pipeline, green `#0e6b0e` for containment and for things that are real, amber `#a86e00` for conditional or simulated, red `#b4283f` for the attack and for hard limits.

The exported SVGs embed their font, so they render identically here, offline, and in a deck. There is nothing external to fetch.

Related

- [Catalog one-pager](#) — what this demo is and who it is for
- [Bill of materials](#) — everything needed to run it
- [Set up at your Hub](#) — branding, join link, checklist